

# Numerical optimization

Mines Nancy – Fall 2024

session 6 – conjugate gradient methods

**Lecturer:** Julien Flamant (CNRS, CRAN)

✉ `julien.flamant@univ-lorraine.fr`

📍 Office 425, FST 1er cycle

**Course material:**

🌐 `arche.univ-lorraine.fr/course/view.php?id=74098`

🐙 `github.com/jflamant/mines-nancy-fall24-optimization`

# Setting

We consider the linear system of  $N$  equations

$$\mathbf{Q}\mathbf{x} = \mathbf{p} \tag{\mathcal{L}}$$

where  $\mathbf{x} \in \mathbb{R}^N$ ,  $\mathbf{Q} \in \mathbb{R}^{N \times N}$  and  $\mathbf{p} \in \mathbb{R}^N$ . We assume that  $\mathbf{Q} > 0$ .

The solution to  $(\mathcal{L})$  is of course  $\mathbf{x}^* = \mathbf{Q}^{-1}\mathbf{p}$ .

## Comments

- for large  $N$ , the cost of matrix inversion makes the direct calculation of  $\mathbf{x}^*$  untractable;
- instead, we seek **iterative methods** to solve  $(\mathcal{L})$ ;
- the specific structure of the problem calls for a tailored solution: this is known as the **linear conjugate gradient method**;
- problem  $(\mathcal{L})$  is closely related to least squares problems.

main reference for this lecture: Nocedal and Wright (2006), Chapter 5.

## Reminder: least squares normal equations

Consider the least squares problem

$$\min_{\mathbf{x} \in \mathbb{R}^N} \|\mathbf{y} - \mathbf{A}\mathbf{x}\|_2^2, \quad \mathbf{A} \in \mathbb{R}^{M \times N}, \mathbf{y} \in \mathbb{R}^M$$

where we assume that  $\text{rank } \mathbf{A} = N$ .

Recall the normal equations for this optimization problem:

$$\nabla f(\mathbf{x}) = \mathbf{0} \Leftrightarrow \mathbf{A}^\top \mathbf{A} \mathbf{x} = \mathbf{A}^\top \mathbf{y}$$

Pose  $\mathbf{Q} = \mathbf{A}^\top \mathbf{A}$  and  $\mathbf{p} = \mathbf{A}^\top \mathbf{y}$ , we obtain the linear system ( $\mathcal{L}$ ). (recall that  $\mathbf{Q} \succ 0$  since  $\text{rank } \mathbf{A} = N$ ). Hence the two problems are equivalent!

The linear conjugate gradient method is effective  
for solving large-scale least squares problems

# Outline

- 1 Linear conjugate gradient method
- 2 Nonlinear conjugate gradient method

# General idea: exploiting the conjugacy property

Conjugate gradient methods rely on the *conjugacy* property.

Let  $\mathbf{Q} \in \mathbb{R}^{N \times N}$  a positive definite matrix.

A set of non-zero vectors  $\{\mathbf{d}_0, \mathbf{d}_1, \dots, \mathbf{d}_\ell\}$  is said to be *conjugate* with respect to  $\mathbf{Q}$  if

$$\mathbf{d}_i^\top \mathbf{Q} \mathbf{d}_j = 0 \text{ for all } i \neq j.$$

Note that this implies that the  $\mathbf{d}_i$ 's are linearly independent.

# General idea: exploiting the conjugacy property

Conjugate gradient methods rely on the *conjugacy* property.

Let  $\mathbf{Q} \in \mathbb{R}^{N \times N}$  a positive definite matrix.

A set of non-zero vectors  $\{\mathbf{d}_0, \mathbf{d}_1, \dots, \mathbf{d}_\ell\}$  is said to be *conjugate* with respect to  $\mathbf{Q}$  if

$$\mathbf{d}_i^T \mathbf{Q} \mathbf{d}_j = 0 \text{ for all } i \neq j.$$

Note that this implies that the  $\mathbf{d}_i$ 's are linearly independent.

## Conjugate direction method

Let  $\mathbf{x}^{(0)} \in \mathbb{R}^N$  a starting point and  $\{\mathbf{d}_0, \mathbf{d}_1, \dots, \mathbf{d}_\ell\}$  a set of conjugate directions.

At iteration  $k$ , generate the new iterate as

$$\mathbf{x}^{(k+1)} = \mathbf{x}^{(k)} + \alpha_k \mathbf{d}_k$$

where  $\alpha_k$  is the optimal step size for the quadratic function associated to the systems of equations  $\mathbf{Q}\mathbf{x} = \mathbf{p}$ .

# Conjugate direction algorithm (I)

**Computation of  $\alpha_k$**  Let  $\phi(\alpha) = \frac{1}{2}\mathbf{t}(\alpha)^\top \mathbf{Q}\mathbf{t}(\alpha) - \mathbf{p}^\top \mathbf{t}(\alpha)$  where  $\mathbf{t}(\alpha) = \mathbf{x}^{(k)} + \alpha \mathbf{d}_k$ . Hence  $\alpha_k = \arg \min_{\alpha} \phi(\alpha)$ .

It is a quadratic convex function of  $\alpha$ , so it suffices to cancel out the derivative to get  $\alpha_k$ .

$$\begin{aligned}\phi'(\alpha) &= \mathbf{d}_k^\top \mathbf{Q} [\mathbf{x}^{(k)} + \alpha \mathbf{d}_k] - \mathbf{p}^\top \mathbf{d}_k \\ &= +\alpha \mathbf{d}_k^\top \mathbf{Q} \mathbf{d}_k + \mathbf{d}_k^\top [\mathbf{Q} \mathbf{x}^{(k)} - \mathbf{p}]\end{aligned}$$

Define  $\mathbf{r}_k = \mathbf{Q}\mathbf{x}^{(k)} - \mathbf{p}$  the residual at iteration  $k$ . Then

$$\alpha_k = -\frac{\mathbf{d}_k^\top \mathbf{r}_k}{\mathbf{d}_k^\top \mathbf{Q} \mathbf{d}_k}$$

## Conjugate direction algorithm (II)

### Theorem (Nocedal and Wright (2006))

*The sequence  $\{\mathbf{x}^{(k+1)}\}$  generated by*

$$\mathbf{x}^{(k+1)} = \mathbf{x}^{(k)} + \alpha_k \mathbf{d}_k, \text{ where } \alpha_k = -\frac{\mathbf{d}_k^\top \mathbf{r}_k}{\mathbf{d}_k^\top \mathbf{Q} \mathbf{d}_k}$$

*where the  $\{\mathbf{d}_k\}$  are conjugate directions converges to the solution  $\mathbf{x}^*$  of the linear system in at most  $N$  steps.*

**Proof** Let us consider  $N$  conjugate directions. This implies that the directions are linearly independent and thus they must span  $\mathbb{R}^N$ . In particular, there exist  $\sigma_0, \dots, \sigma_{N-1} \in \mathbb{R}$  such that

$$\mathbf{x}^* - \mathbf{x}^{(0)} = \sum_{i=0}^{N-1} \sigma_i \mathbf{d}_i$$



## Proof (continued)

By using the conjugacy property, we can obtain the value of  $\sigma_k$  as

$$\mathbf{d}_k^\top \mathbf{Q} [\mathbf{x}^* - \mathbf{x}^{(0)}] = \mathbf{d}_k^\top \mathbf{Q} \left[ \sum_{i=0}^{N-1} \sigma_i \mathbf{d}_i \right] = \sigma_k \mathbf{d}_k^\top \mathbf{Q} \mathbf{d}_k$$

hence  $\sigma_k = \frac{\mathbf{d}_k^\top \mathbf{Q} [\mathbf{x}^* - \mathbf{x}^{(0)}]}{\mathbf{d}_k^\top \mathbf{Q} \mathbf{d}_k}$ . Next we prove that  $\alpha_k = \sigma_k$ .

By the recursion formula, we have for  $k > 0$ ,  $\mathbf{x}^{(k)} = \mathbf{x}^{(0)} + \sum_{i=0}^{k-1} \alpha_i \mathbf{d}_i$ .  
Then by conjugacy

$$\mathbf{d}_k^\top \mathbf{Q} \mathbf{x}^{(k)} = \mathbf{d}_k^\top \mathbf{Q} \mathbf{x}^{(0)} \Leftrightarrow \mathbf{d}_k^\top \mathbf{Q} [\mathbf{x}^{(k)} - \mathbf{x}^{(0)}] = 0$$

and as a result

$$\mathbf{d}_k^\top \mathbf{Q} [\mathbf{x}^* - \mathbf{x}^{(0)}] = \mathbf{d}_k^\top \mathbf{Q} [\mathbf{x}^* - \mathbf{x}^{(k)}] = \mathbf{d}_k^\top [\mathbf{p} - \mathbf{Q} \mathbf{x}^{(k)}] = -\mathbf{d}_k^\top \mathbf{r}_k$$

which permits to conclude that  $\alpha_k = \sigma_k$ .

# Additional properties

## Properties of the conjugate direction method

Let  $\{\mathbf{x}^{(k)}\}$  be generated by the conjugate direction method. Then, for an iteration  $k$ , one has

- $\mathbf{r}_k = \mathbf{r}_{k-1} + \alpha_{k-1} \mathbf{Q} \mathbf{d}_{k-1}$
- $\mathbf{d}_i^\top \mathbf{r}_k = 0$  for  $i = 0, 1, \dots, k-1$
- $\mathbf{x}^{(k)}$  is the minimizer of  $\phi(\mathbf{x}) = \frac{1}{2} \mathbf{x}^\top \mathbf{Q} \mathbf{x} - \mathbf{p}^\top \mathbf{x}$  over the set  $\{\mathbf{x} \mid \mathbf{x}^{(0)} + \text{span}\{\mathbf{d}_0, \mathbf{d}_1, \dots, \mathbf{d}_{k-1}\}\}$

## Comments

- See Nocedal and Wright (2006), p. 106 for a proof.
- The residual  $\mathbf{r}_k$  only depends on the previous one and on the current iteration ( $\alpha_{k-1}, \mathbf{d}_{k-1}$ )
- the residual  $\mathbf{r}_k$  is orthogonal to the  $k$  previous conjugate directions

# The linear conjugate gradient method

The conjugate directions method requires a set of conjugate directions  $\{\mathbf{d}_k\}$ , which can be determined in advance using, e.g.,

- eigenvalue decomposition of  $\mathbf{Q}$ ;
- modification of the Gram-Schmidt process to produce a set of conjugate directions

→ too costly for large scale applications

Solution: (linear) conjugate gradient (CG) method

- directions  $\mathbf{d}_k$  are computed iteratively;
- each new direction  $\mathbf{d}_k$  use only  $\mathbf{d}_{k-1}$ ; it is automatically conjugated to all the previous directions  $\mathbf{d}_0, \mathbf{d}_1, \dots, \mathbf{d}_{k-1}$
- cheap computational cost and memory storage.

# The linear conjugate gradient method

**Idea:** choose the direction  $\mathbf{d}_k$  as a linear combination of  $-\mathbf{r}_k$  and  $\mathbf{d}_{k-1}$

$$\mathbf{d}_k = -\mathbf{r}_k + \beta_k \mathbf{d}_{k-1}$$

where  $\beta_k$  is set to impose conjugation between  $\mathbf{d}_k$  and  $\mathbf{d}_{k-1}$ :

$$\beta_k = \frac{\mathbf{d}_{k-1}^\top \mathbf{Q} \mathbf{r}_k}{\mathbf{d}_{k-1}^\top \mathbf{Q} \mathbf{d}_{k-1}}$$

This gives us a preliminary version of CG.

Starting from  $\mathbf{x}^{(0)} \in \mathbb{R}^N$  and  $\mathbf{d}_0 = -\mathbf{r}_0$ , iterate until convergence

$$\mathbf{x}^{(k+1)} = \mathbf{x}^{(k)} + \alpha_k \mathbf{d}_k, \text{ where } \alpha_k = -\frac{\mathbf{d}_k^\top \mathbf{r}_k}{\mathbf{d}_k^\top \mathbf{Q} \mathbf{d}_k}$$

$$\mathbf{d}_{k+1} = -\mathbf{r}_{k+1} + \beta_{k+1} \mathbf{d}_k, \text{ where } \beta_{k+1} = \frac{\mathbf{d}_k^\top \mathbf{Q} \mathbf{r}_{k+1}}{\mathbf{d}_k^\top \mathbf{Q} \mathbf{d}_k}$$

## Results for CG: theorem

This first version is practical for *studying properties* of CG.

Krylov subspace: useful tool in numerical linear algebra

$$\mathcal{K}(\mathbf{r}; k) = \text{span}\{\mathbf{r}, \mathbf{Q}\mathbf{r}, \mathbf{Q}^2\mathbf{r}, \dots, \mathbf{Q}^k\mathbf{r}\}$$

### Theorem (Nocedal and Wright (2006))

*Consider the  $k$ -th iteration of the CG method (not the final one). Then*

$$\mathbf{r}_k^\top \mathbf{r}_i = 0 \text{ for } i = 0, 1, \dots, k-1$$

$$\text{span}\{\mathbf{r}_0, \mathbf{r}_1, \dots, \mathbf{r}_k\} = \mathcal{K}(\mathbf{r}_0; k)$$

$$\text{span}\{\mathbf{d}_0, \mathbf{d}_1, \dots, \mathbf{d}_k\} = \mathcal{K}(\mathbf{r}_0; k)$$

$$\mathbf{d}_k^\top \mathbf{Q} \mathbf{d}_i = 0 \text{ for } i = 0, 1, \dots, k-1$$

*and thus the sequence  $\{\mathbf{x}^{(k)}\}$  converges to  $\mathbf{x}^*$  in at most  $N$  iterations.*

**proof:** see Nocedal and Wright (2006), pp. 109-111.

# Results for CG: some comments

## Comments on the theorem

- residuals  $\{\mathbf{r}_k\}$  are mutually orthogonal;
- Residuals  $\mathbf{r}_k$  and search directions  $\mathbf{d}_k$  all belong to the Krylov subspace of order  $k$  associated with  $\mathbf{r}_0$ ;
- the search directions  $\mathbf{d}_0, \mathbf{d}_1, \dots, \mathbf{d}_{N-1}$  are conjugate wrt  $\mathbf{Q}$
- by the theorem on conjugate directions (slide 6), this implies termination in at most  $N$  steps.
- these results all depend on the choice of  $\mathbf{d}_0$ ! In fact, if one chooses  $\mathbf{d}_0 \neq -\mathbf{r}_0$  (i.e., different from the steepest direction at  $\mathbf{x}^{(0)}$ ) then the theorem does not longer holds.

# Standard CG algorithm

Using the Theorems on slide 8 and 11, we can improve the computations of  $\alpha_k$  and  $\beta_{k+1}$  in CG:

$$\alpha_k = -\frac{\mathbf{d}_k^\top \mathbf{r}_k}{\mathbf{d}_k^\top \mathbf{Q} \mathbf{d}_k} = -\frac{[-\mathbf{r}_k + \beta_k \mathbf{d}_{k-1}]^\top \mathbf{r}_k}{\mathbf{d}_k^\top \mathbf{Q} \mathbf{d}_k} = \frac{\|\mathbf{r}_k\|_2^2}{\mathbf{d}_k^\top \mathbf{Q} \mathbf{d}_k}$$

Moreover, observe that  $\mathbf{r}_{k+1} - \mathbf{r}_k = \alpha_k \mathbf{Q} \mathbf{d}_k$  so that

$$\beta_{k+1} = \frac{\mathbf{d}_k^\top \mathbf{Q} \mathbf{r}_{k+1}}{\mathbf{d}_k^\top \mathbf{Q} \mathbf{d}_k} = \frac{\mathbf{d}_k^\top \mathbf{Q} \mathbf{d}_k}{\|\mathbf{r}_k\|_2^2} \frac{[\mathbf{r}_{k+1} - \mathbf{r}_k]^\top \mathbf{r}_{k+1}}{\mathbf{d}_k^\top \mathbf{Q} \mathbf{d}_k} = \frac{\|\mathbf{r}_{k+1}\|_2^2}{\|\mathbf{r}_k\|_2^2}$$

where we used the orthogonality of residuals.

# Standard CG algorithm

## For which problems?

- Solve "square" linear systems of the form  $\mathbf{Q}\mathbf{x} = \mathbf{p}$  with  $\mathbf{Q} > 0$
- For strictly convex quadratic problems  $\min_{\mathbf{x}} \frac{1}{2}\mathbf{x}^\top \mathbf{Q}\mathbf{x} - \mathbf{p}^\top \mathbf{x}$   
(or full column-rank least-squares problems)

## Standard linear CG algorithm

- 1: **input:**  $\mathbf{x}^{(0)} \in \mathbb{R}^N$
- 2: Set  $\mathbf{r}_0 = \mathbf{Q}\mathbf{x}^{(0)} - \mathbf{p}$  and  $\mathbf{d}_0 = -\mathbf{r}_0$ .
- 3:  $k = 0$
- 4: **while**  $\|\mathbf{r}_k\|_2 \neq 0$  **do**
- 5:    $\mathbf{x}^{(k+1)} = \mathbf{x}^{(k)} + \alpha_k \mathbf{d}_k$ , where  $\alpha_k = \frac{\|\mathbf{r}_k\|_2^2}{\mathbf{d}_k^\top \mathbf{Q} \mathbf{d}_k}$
- 6:    $\mathbf{d}_{k+1} = -\mathbf{r}_{k+1} + \beta_{k+1} \mathbf{d}_k$ , where  $\beta_{k+1} = \frac{\|\mathbf{r}_{k+1}\|_2^2}{\|\mathbf{r}_k\|_2^2}$
- 7:    $k = k + 1$
- 8: **end while**
- 9: **return**  $\mathbf{x}^{(k)}$



# Numerical performance and conditioning

- convergence to  $\mathbf{x}^*$  is guaranteed after at most  $N$  iterations
- but it can be much faster!

**Trivial example** Solve system  $\mathbf{Q}\mathbf{x} = \mathbf{p}$  using CG when  $\mathbf{Q} = \lambda\mathbf{I}_N$ .

first residual  $\mathbf{r}_0 = \lambda\mathbf{x}^{(0)} - \mathbf{p}$  and direction  $\mathbf{d}_0 = -\mathbf{r}_0$ .

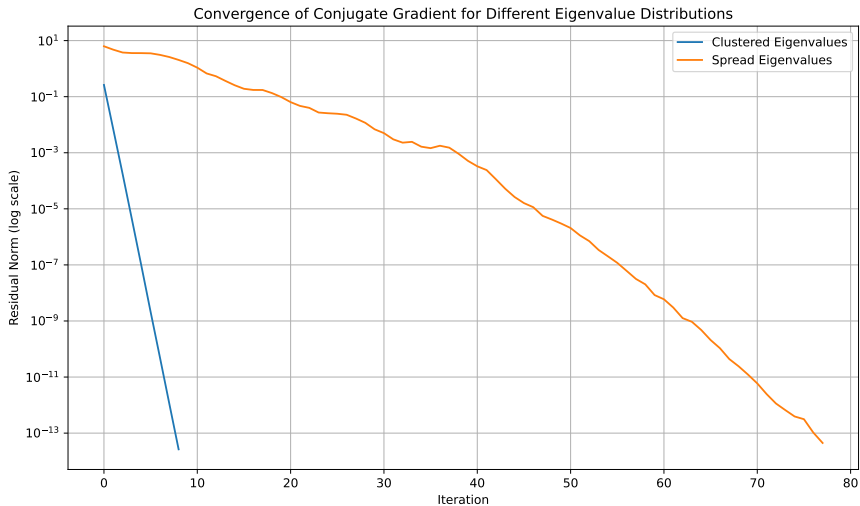
first iteration:

$$\mathbf{x}^{(1)} = \mathbf{x}^{(0)} - \frac{\|\lambda\mathbf{x}^{(0)} - \mathbf{p}\|^2}{\lambda\|\lambda\mathbf{x}^{(0)} - \mathbf{p}\|^2} \lambda(\lambda\mathbf{x}^{(0)} - \mathbf{p}) = \mathbf{x}^{(0)} - \mathbf{x}^{(0)} + \lambda^{-1}\mathbf{p} = \mathbf{x}^*$$

Then  $\mathbf{r}_{k+1} = \mathbf{0}$  and  $\mathbf{d}_{k+1} = \mathbf{0}$ . CG stops after exactly 1 iteration.

- convergence usually depends on the **distribution of eigenvalues of  $\mathbf{Q}$**
- it can be shown that if  $\mathbf{Q}$  has  $r$  distinct eigenvalues, then CG converges in at most  $r$  iterations.

# Numerical experiment for $N = 100$



# Summary for linear conjugate gradient

- exact convergence guarantees in at most  $N$  iterations  
compare with results of session 5 on descent methods
- state of the art for solving linear least squares in high-dimensions
- the closer the condition number of  $\mathbf{Q}$  is to 1, the faster the convergence  
→ use of preconditioners to improve the condition number of  $\mathbf{Q}$
- implemented in many scientific libraries, e.g.,  
`scipy.sparse.linalg.cg` (Python) or `cgs` (Matlab)

# Outline

- 1 Linear conjugate gradient method
- 2 Nonlinear conjugate gradient method**

# Motivation

Linear conjugate gradient solves linear system  $\mathbf{Q}\mathbf{x} = \mathbf{p}$ , or equivalently,

$$\min_{\mathbf{x} \in \mathbb{R}^N} f(\mathbf{x}) := \frac{1}{2} \mathbf{x}^\top \mathbf{Q} \mathbf{x} - \mathbf{p}^\top \mathbf{x}, \quad \mathbf{Q} > 0$$

**Nonlinear** conjugate gradient methods extend CG to non-quadratic objective functions  $f$ .

Several variants exists, two popular ones:

- the Fletcher-Reeves method
- the Polak-Ribière method

the key principle is the following:

replace the residual  $\mathbf{r}_k$  (which is the gradient of quadratic  $f$ ) by the evaluation of the gradient of the nonlinear objective  $f$  at iteration  $\mathbf{x}^{(k)}$ .

# Fletcher-Reeves vs Polyak-Ribière

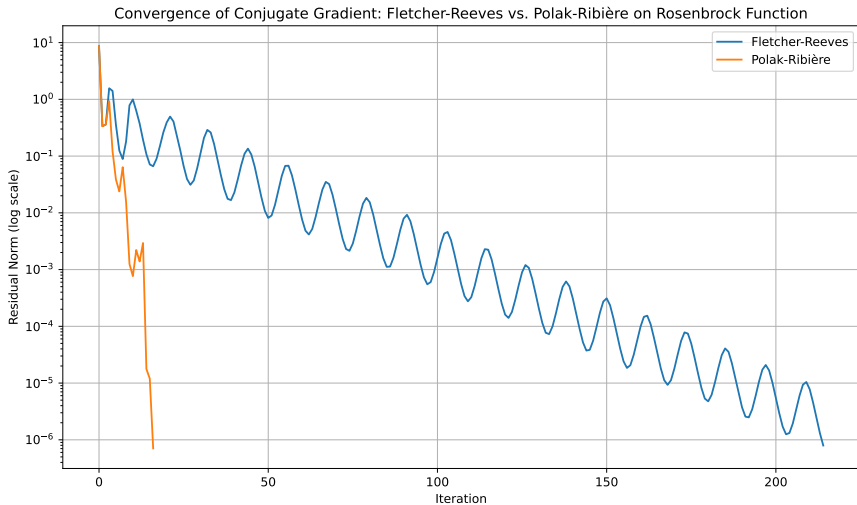
## Fletcher-Reeves-CG algorithm

```
1: input:  $\mathbf{x}^{(0)} \in \mathbb{R}^N$ 
2: Set  $\nabla f_0 = \nabla f(\mathbf{x}^{(0)})$  and set  $\mathbf{d}_0 = -\nabla f_0$ .
3:  $k = 0$ 
4: while  $\|\nabla f_k\|_2 \neq 0$  do
5:    $\mathbf{x}^{(k+1)} = \mathbf{x}^{(k)} + \alpha_k \mathbf{d}_k$ , where  $\alpha_k$  is computed by line search
6:    $\mathbf{d}_{k+1} = -\nabla f_{k+1} + \beta_{k+1}^{\text{FR}} \mathbf{d}_k$ , where  $\beta_{k+1}^{\text{FR}} = \frac{\|\nabla f_{k+1}\|_2^2}{\|\nabla f_k\|_2^2}$ 
7:    $k = k + 1$ 
8: end while
9: return  $\mathbf{x}^{(k)}$ 
```

- $\alpha_k$  must be chosen with care, to that  $\mathbf{d}_{k+1}$  remains a descent direction.  $\rightarrow$  use of strong Wolfe conditions
- Polak-Ribière: replace  $\beta_{k+1}^{\text{FR}}$  by

$$\beta_{k+1}^{\text{PR}} = \frac{\nabla f_{k+1}^\top (\nabla f_{k+1} - \nabla f_k)}{\|\nabla f_k\|_2^2} \quad \text{or} \quad \beta_{k+1}^+ = \max(0, \beta_{k+1}^{\text{PR}}) \quad (\text{more robust})$$

# Numerical experiment on Rosenbrock function



# Summary

- PR is usually more efficient than FR; but this not a general rule and depends on the context
- many other variants exist: Hestenes-Stiefel, Dai-Yuan, etc.
- convergence analysis is still possible, but much more technical than for the linear conjugate gradient method: see Nocedal and Wright (2006), Chapter 5 for details.
- the Polak-Ribière method is implemented in Python:  
check out `scipy.optimize.fmin_cg`