

Numerical optimization

Mines Nancy – Fall 2024

session 8 – constrained convex optimization II
Algorithms

Lecturer: Julien Flamant (CNRS, CRAN)

✉ julien.flamant@univ-lorraine.fr

📍 Office 425, FST 1er cycle

Course material:

🌐 arche.univ-lorraine.fr/course/view.php?id=74098

🐙 github.com/jflamant/mines-nancy-fall24-optimization

Setting: constrained convex optimization

We consider optimization problems of the form

$$\min_{\mathbf{x} \in \Omega} f_0(\mathbf{x}) \quad (\mathcal{P})$$

where we make the following assumptions:

- the objective $f_0 : \mathcal{D}_0 \subseteq \mathbb{R}^N \rightarrow \mathbb{R}$ is convex on its domain \mathcal{D}_0
- the set of constraints $\Omega \subset \mathbb{R}^N$ is convex. We assume that Ω is described by M inequality constraints and P equality constraints such that

$$\Omega = \left\{ \mathbf{x} \in \mathbb{R}^N \left| \begin{array}{ll} f_i(\mathbf{x}) \leq 0 & , \quad i = 1, \dots, M \\ \mathbf{a}_i^\top \mathbf{x} = b_i & , \quad i = 1, \dots, P \end{array} \right. \right\}$$

where the functions $f_i : \mathbb{R}^N \rightarrow \mathbb{R}$ are convex, hence Ω is a convex set. (prove it)

(\mathcal{P}) is a constrained convex optimization problem

Agenda

Session 7 - 03/12/24: Theory

- Characterization of solutions
- Lagrangian duality
- Karush-Kuhn-Tucker (KKT) conditions

Session 8 - 10/12/24: Algorithms

- Penalty methods
- Projected gradient
- Dual methods: Uzawa's algorithm
- Disciplined convex programming with CVXPY

Main reference

Convex Optimization, Boyd and Vandenberghe (2004) - Chapters 4, 5

Outline

- ① Penalty method
- ② Projected gradient
- ③ Dual methods: Uzawa's algorithm
- ④ Disciplined Convex programming with CVXPY

Penalty method principle

Consider the optimization problem

$$\min_{\mathbf{x} \in \Omega} f_0(\mathbf{x}) \quad (\mathcal{P})$$

where f_0 is the objective and $\Omega \subset \mathbb{R}^N$ is the constraint set.

Penalty method

Replace (\mathcal{P}) with the unconstrained optimization problem

$$\min_{\mathbf{x} \in \mathbb{R}^N} f_0(\mathbf{x}) + \rho p(\mathbf{x}), \quad \rho > 0 \quad (\mathcal{P}_\rho)$$

where $p: \mathbb{R}^N \rightarrow \mathbb{R}$ is the **penalty function** such that

- p is continuous
- $p(\mathbf{x}) > 0$ if $\mathbf{x} \notin \Omega$
- $p(\mathbf{x}) = 0$ if $\mathbf{x} \in \Omega$.

The **penalty parameter** ρ controls the cost of violating the constraints.

Choosing penalty functions

Recall Ω is often parameterized by equality and inequality constraints.

$$\Omega = \left\{ \mathbf{x} \in \mathbb{R}^N \left| \begin{array}{ll} f_i(\mathbf{x}) \leq 0 & , \quad i = 1, \dots, M \\ h_i(\mathbf{x}) = 0 & , \quad i = 1, \dots, P \end{array} \right. \right\}$$

Typical choices of penalty functions: use **quadratic functions**

- **inequality constraint** $f_i(\mathbf{x}) \leq 0$: take

$$p(\mathbf{x}) = \frac{1}{2} [\max(0, f_i(\mathbf{x}))]^2 := \frac{1}{2} [f_i^+(\mathbf{x})]^2$$

- **equality constraint** $h_i(\mathbf{x}) = 0$: take

$$p(\mathbf{x}) = \frac{1}{2} (h_i(\mathbf{x}))^2$$

for multiple constraints, simply sum all penalties $p(\mathbf{x}) = \sum_i p_i(\mathbf{x})$

Remark: can consider also non-smooth penalties such as ℓ_1 with very interesting properties (outside the scope of this lecture).

Penalty method algorithm

$$\min_{\mathbf{x} \in \Omega} f_0(\mathbf{x}) \quad (\mathcal{P})$$

Solve (\mathcal{P}) by a succession of penalized problems

$$\min_{\mathbf{x} \in \mathbb{R}^N} f_0(\mathbf{x}) + \rho_k p(\mathbf{x}), \quad \rho_k > 0 \quad (\mathcal{P}_{\rho_k})$$

with increasing penalties $\rho_1 < \rho_2 < \dots < \rho_k \rightarrow +\infty$

Theorem (Convergence of the penalty method)

Let f_0 and p be continuous. Let $\mathbf{x}^{(k)}$ be a minimizer of \mathcal{P}_{ρ_k} . Suppose the sequence $\{\rho_k\}$ is strictly increasing with $\rho_k \rightarrow \infty$ as $k \rightarrow \infty$. Then $\mathbf{x}^{(k)} \rightarrow \mathbf{x}^$, where \mathbf{x}^* is a solution of (\mathcal{P}) .*

Computing solutions of penalized problems

Consider a problem with one inequality constraint $f_1(\mathbf{x}) \leq 0$.

The quadratic penalty is $p(\mathbf{x}) = \frac{1}{2} [\max(0, f_1(\mathbf{x}))]^2 := \frac{1}{2} [f_1^+(\mathbf{x})]^2$

How to compute the gradient?

The main issue is that f_1^+ is not differentiable at points \mathbf{x} such that $f_1(\mathbf{x}) = 0$.

However, remark that p is a function of f_1^+ only s.t. $p(\mathbf{x}) = \gamma(f_1^+(\mathbf{x}))$, where γ is the quadratic function. In particular, $\frac{\partial \gamma}{\partial y}(0) = 0$, which implies that $p(\mathbf{x})$ is differentiable whenever f_1^+ is.

Using the chain rule

$$\nabla p(\mathbf{x}) = f_1^+(\mathbf{x}) \nabla f_1(\mathbf{x}) = \begin{cases} f_1(\mathbf{x}) \nabla f_1(\mathbf{x}) & \text{if } f_1(\mathbf{x}) \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

Example

Using the penalty method, solve the following problem:

$$\begin{array}{ll}\text{minimize} & x \\ \text{subject to} & x \geq 1\end{array}$$

Example

Using the penalty method, solve the following problem:

$$\begin{array}{ll}\text{minimize} & x \\ \text{subject to} & x \geq 1\end{array}$$

Solution Write the penalized problem:

$$\text{minimize} \quad x + \frac{\rho}{2}(\max(0, 1 - x))^2 := f_{\rho}(x), \quad \rho > 0$$

The objective is convex, with gradient $f'_{\rho}(x) = 1 - \rho \max(0, 1 - x)$.
Solving $f'_{\rho}(x) = 0$ gives $x = 1 - 1/\rho$, which converges to $x = 1$ as $\rho \rightarrow \infty$.

Another example

Using the penalty method, solve the following problem:

$$\begin{array}{ll}\text{minimize} & \frac{1}{2} \|\mathbf{x}\|_2^2 \\ \text{subject to} & \mathbf{Ax} = \mathbf{b}\end{array}$$

Solution Write the penalized problem:

$$\text{minimize} \quad \frac{1}{2} \|\mathbf{x}\|^2 + \frac{\rho}{2} \|\mathbf{Ax} - \mathbf{b}\|_2^2, \quad \rho > 0$$

The objective is convex, with gradient $\nabla f_\rho(\mathbf{x}) = \mathbf{x} + \rho \mathbf{A}^\top (\mathbf{Ax} - \mathbf{b})$. Solving $\nabla f_\rho(\mathbf{x}) = 0$ gives $\mathbf{x}_\rho = \rho(\mathbf{I} + \rho \mathbf{A}^\top \mathbf{A})^{-1} \mathbf{A}^\top \mathbf{b}$.

Tedious calculations (e.g., using the SVD) show that, in the limit $\rho \rightarrow \infty$, $\mathbf{x}_\rho \rightarrow \mathbf{x}^*$ with $\mathbf{x}^* = \mathbf{A}^\top (\mathbf{AA}^\top)^{-1} \mathbf{b}$.

Practical algorithm

$$\min_{\mathbf{x} \in \Omega} f_0(\mathbf{x}) \quad (\mathcal{P})$$

Penalty algorithm

Choose a penalty function $p(\mathbf{x})$ for the problem (\mathcal{P}) , e.g., a quadratic penalty: $p(\mathbf{x}) = \sum_{i=1}^M [f_i^+(\mathbf{x})]^2 + \sum_{i=1}^P h_i^2(\mathbf{x})$.

Choose a stopping criterion. Then the algorithm reads:

- 1: Fix increase parameter $s > 1$ and initial penalty $\rho_0 > 0$.
- 2: $k := 0$
- 3: **while** stopping criterion is not satisfied **do**
- 4: $\mathbf{x}^{(k+1)} := \arg \min (f_0(\mathbf{x}) + \rho_k p(\mathbf{x}))$ (possibly using $\mathbf{x}^{(k)}$)
- 5: $k := k + 1$
- 6: $\rho_{k+1} := s\rho_k$
- 7: **end while**
- 8: **return** (approximate) solution $\mathbf{x}^{(k)}$

Outline

- ① Penalty method
- ② Projected gradient**
- ③ Dual methods: Uzawa's algorithm
- ④ Disciplined Convex programming with CVXPY

Motivation

Consider the general constrained optimization problem

$$\min_{\mathbf{x} \in \Omega} f_0(\mathbf{x}) \quad (\mathcal{P})$$

goal: design a simple algorithm to solve (\mathcal{P})

recall from lecture 4: for **unconstrained problems** ($\Omega = \mathbb{R}^N$) with differentiable objective, one simple algorithm is **gradient descent** (GD):

$$\mathbf{x}^{(k+1)} = \mathbf{x}^{(k)} - \alpha_k \nabla f(\mathbf{x}^{(k)}), \quad \alpha_k > 0, \quad k = 0, 1, \dots$$

for **constrained problems**, a natural generalization is **projected gradient descent** (PGD)

combines **GD** with (Euclidean) **projection onto Ω**

Projection onto a set

The (Euclidean) projection operator $P_\Omega : \mathbb{R}^N \rightarrow \mathbb{R}^N$ onto Ω is defined as

$$P_\Omega(\mathbf{x}_0) = \arg \min_{\mathbf{x} \in \Omega} \|\mathbf{x} - \mathbf{x}_0\|_2^2$$

- computing $P_\Omega(\mathbf{x}_0)$ is itself an optimization problem
- $P_\Omega(\mathbf{x}_0)$ returns the closest point(s) (in Euclidean norm) to \mathbf{x}_0 that belongs to the set Ω
- when Ω is closed and convex the projection is unique
- if $\mathbf{x}_0 \in \Omega$, then $P_\Omega(\mathbf{x}_0) = \mathbf{x}_0$
- in many important cases, $P_\Omega(\mathbf{x}_0)$ can be computed explicitly

$$P_{\mathbb{R}_+^N}(\mathbf{x}_0) = [\max(0, [\mathbf{x}_0]_i)]_{i=1}^N \quad (\text{entry-wise maximum})$$

exercise: compute the projection onto the unit ball in \mathbb{R}^N ,
 $\Omega = \{\mathbf{x} \mid \|\mathbf{x}\| \leq 1\}$

Projected gradient method

$$\min_{\mathbf{x} \in \Omega} f_0(\mathbf{x}) \quad (\mathcal{P})$$

Projected gradient algorithm


- 1: starting point $\mathbf{x}^{(0)} \in \mathbb{R}^N$
- 2: $k := 0$
- 3: **while** stopping criterion is not satisfied **do**
- 4: $\mathbf{x}^{(k+1)} := P_{\Omega}(\mathbf{x}^{(k)} - \alpha_k \nabla f_0(\mathbf{x}^{(k)}))$
- 5: $k := k + 1$
- 6: **end while**
- 7: **return** solution $\mathbf{x}^{(k)}$

typical step-size strategies:

- fixed step-size $\alpha_k = \alpha$ for all k
- use backtracking strategies to determine α_k at each iteration

Example

$$\begin{aligned} &\text{minimize} && (x_1 - 2)^2 + (x_2 - 1)^2 \\ &\text{subject to} && x_1^2 - x_2 \leq 0 \\ &&& x_1 + x_2 \leq 2 \end{aligned}$$



Comments on projected gradient descent

- PGD is **particularly interesting** when the projection operator is either **explicit, easy to solve, or computationally efficient**.
- Theoretical analysis of PGD is similar to that of GD (see session 5). In particular, if f_0 is convex and L -smooth, then convergence is in $\mathcal{O}(1/K)$ (like GD)
- PGD is a special case of **proximal gradient method**, used to solve problems of the form $\min f(\mathbf{x}) + g(\mathbf{x})$, where f is differentiable and g is not. To recover PGD, take g to be the indicator function on Ω
- sufficient decrease condition in backtracking should be adapted to account for the projection step.

Outline

- ① Penalty method
- ② Projected gradient
- ③ Dual methods: Uzawa's algorithm**
- ④ Disciplined Convex programming with CVXPY

Reminder: primal and dual problems

Primal problem

$$\begin{array}{ll}\text{minimize} & f_0(\mathbf{x}) \\ \text{subject to} & f_i(\mathbf{x}) \leq 0, \quad i = 1, \dots, M \\ & h_i(\mathbf{x}) = 0, \quad i = 1, \dots, P\end{array} \quad (\mathcal{P})$$

Dual problem

$$\begin{array}{ll}\text{maximize} & g(\boldsymbol{\lambda}, \boldsymbol{\nu}) := \inf_{\mathbf{x}} L(\mathbf{x}, \boldsymbol{\lambda}, \boldsymbol{\nu}) \\ \text{subject to} & \boldsymbol{\lambda} \geq 0\end{array} \quad (\mathcal{D})$$

where the L is the Lagrangian

$$L(\mathbf{x}, \boldsymbol{\lambda}, \boldsymbol{\nu}) = f_0(\mathbf{x}) + \sum_{i=1}^M \lambda_i f_i(\mathbf{x}) + \sum_{i=1}^P \nu_i h_i(\mathbf{x})$$

Principle of Uzawa's algorithm

- can be viewed as **dual projected gradient**
→ simple projection of λ onto \mathbb{R}_+^P :
(entrywise) $P_{\mathbb{R}_+^P}(\lambda) = \max(0, \lambda)$
- dual is a maximization problem → **gradient ascent**
- when it converges, iterates tends to a saddle point of L (therefore algorithm requires strong duality)

Uzawa's algorithm

- 1: initial Lagrange multipliers $\lambda^{(0)} \geq 0$ and $\nu^{(0)}$.
- 2: $k := 0$
- 3: **while** stopping criterion is not satisfied **do**
- 4: $\mathbf{x}^{(k+1)} := \arg \min_{\mathbf{x}} L(\mathbf{x}, \lambda^{(k)}, \nu^{(k)})$
- 5: $\lambda^{(k+1)} = P_{\mathbb{R}_+^P} \left(\lambda^{(k)} + \alpha_k [f_1(\mathbf{x}^{(k+1)}) \dots f_p(\mathbf{x}^{(k+1)})]^\top \right), \quad \alpha_k > 0 \text{ stepsize}$
- 6: $\nu^{(k+1)} = \nu^{(k)} + \beta_k [h_1(\mathbf{x}^{(k+1)}) \dots h_P(\mathbf{x}^{(k+1)})]^\top, \quad \beta_k > 0 \text{ stepsize}$
- 7: **end while**
- 8: **return** (approximate) solution $\mathbf{x}^{(k)}$

Example

Write Uzawa's iterates for the following problem

$$\begin{array}{ll}\text{minimize} & \frac{1}{2} \|\mathbf{y} - \mathbf{Ax}\|_2^2 \\ \text{subject to} & \mathbf{x} \geq 0\end{array}$$

where \mathbf{A} is full column rank.

Example

Write Uzawa's iterates for the following problem

$$\begin{aligned} & \text{minimize} && \frac{1}{2} \|\mathbf{y} - \mathbf{Ax}\|_2^2 \\ & \text{subject to} && \mathbf{x} \geq 0 \end{aligned}$$

where \mathbf{A} is full column rank.

solution The Lagrangian for the problem is $L(\mathbf{x}, \boldsymbol{\lambda}) = \frac{1}{2} \|\mathbf{y} - \mathbf{Ax}\|_2^2 - \boldsymbol{\lambda}^\top \mathbf{x}$ with gradient with respect to \mathbf{x} given by $\nabla_{\mathbf{x}} L(\mathbf{x}, \boldsymbol{\lambda}) = \mathbf{A}^\top (\mathbf{Ax} - \mathbf{y}) - \boldsymbol{\lambda}$. Therefore we have the iterations

$$\begin{aligned} \mathbf{x}^{(k+1)} &= (\mathbf{A}^\top \mathbf{A})^{-1} [\mathbf{A}^\top \mathbf{y} + \boldsymbol{\lambda}^{(k)}] \\ \boldsymbol{\lambda}^{(k+1)} &= [\boldsymbol{\lambda}^{(k)} - \alpha_k \mathbf{x}^{(k+1)}]^+ \end{aligned}$$

until primal and dual convergence.

Outline

- ① Penalty method
- ② Projected gradient
- ③ Dual methods: Uzawa's algorithm
- ④ Disciplined Convex programming with CVXPY

Disciplined Convex Programming (DCP)

What is it?

A set of rules, a “language” that permits to formulate and implement efficiently convex optimization problems.

→ underlying language of numerical solvers, such as CVXPY (Python) or CVX (Matlab)

CVXPY: Python-based solver using DCP

see <https://www.cvxpy.org/> for tutorial and documentation

check also the JMLR paper:

CVXPY: A Python-Embedded Modeling Language for Convex Optimization, by Diamond and Boyd

CVXPY example [from tutorial]

```
1  import cvxpy as cp
2
3  # Create two scalar optimization variables.
4  x = cp.Variable()
5  y = cp.Variable()
6
7  # Create two constraints.
8  constraints = [x + y == 1,
9                x - y >= 1]
10
11 # Form objective.
12 obj = cp.Minimize((x - y)**2)
13
14 # Form and solve problem.
15 prob = cp.Problem(obj, constraints)
16 prob.solve() # Returns the optimal value.
17 print("status:", prob.status)
18 print("optimal value", prob.value)
19 print("optimal var", x.value, y.value)
```

Exercise

Solve the following problem using CVXPY

$$\begin{array}{ll}\text{minimize} & (x_1 - 2)^2 + (x_2 - 1)^2 \\ \text{subject to} & x_1^2 - x_2 \leq 0 \\ & x_1 + x_2 \leq 2\end{array}$$